

xgboost-for-sales-forecasting

April 22, 2024

1 XGBoost

50 2013 ~2017 2018
One-Hot-Encoding

```
[ ]: import xgboost as xgb
import seaborn as sns
import numpy as np # linear algebra
import matplotlib.pyplot as plt
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from matplotlib.pyplot import figure

[ ]: train_data = pd.read_csv("/kaggle/input/demand-forecasting-kernels-only/train.
    ↪csv")
test_data = pd.read_csv("/kaggle/input/demand-forecasting-kernels-only/test.
    ↪csv")
df = pd.concat([train_data, test_data], sort=False)
df.tail(10)
```

```
[ ]:
      date  store  item  sales      id
44990 2018-03-22    10    50    NaN  44990.0
44991 2018-03-23    10    50    NaN  44991.0
44992 2018-03-24    10    50    NaN  44992.0
44993 2018-03-25    10    50    NaN  44993.0
44994 2018-03-26    10    50    NaN  44994.0
44995 2018-03-27    10    50    NaN  44995.0
44996 2018-03-28    10    50    NaN  44996.0
44997 2018-03-29    10    50    NaN  44997.0
44998 2018-03-30    10    50    NaN  44998.0
44999 2018-03-31    10    50    NaN  44999.0
```

2

2.1 Descriptive Statistics

```
[ ]: #
statistic_sheet_s = df.groupby(["store"]).agg({"sales": ["count", "sum", "mean", "median", "std", "min", "max"]})
statistic_sheet_s.head(2)
```


```
[ ]:      sales
      count      sum      mean median      std  min   max
store
1      91300  4315603.0  47.268379   44.0  24.006252   1.0  155.0
2      91300  6120128.0  67.033165   62.0  33.595810   3.0  231.0
```

```
[ ]: #
statistic_sheet_i = df.groupby(["item"]).agg({"sales": ["count", "sum", "mean", "median", "std", "min", "max"]})
statistic_sheet_i.head(2)
```

```
[ ]:      sales
      count      sum      mean median      std  min   max
item
1      18260  401384.0  21.981599   21.0   8.468922   1.0   59.0
2      18260 1069564.0  58.574151   56.0  20.093015   9.0  150.0
```


```
[ ]: #
stores_sum = df.groupby(["store"], as_index=False).agg({"sales": "sum"}).
    sort_values(by="sales", ascending=False)
sns.barplot(data=stores_sum, x='store', y='sales', color="green", order=stores_sum.
    sort_values('sales').store)
```

```
[ ]: <AxesSubplot: xlabel='store', ylabel='sales'>
```



```
[ ]: #
figure(figsize=(18, 4), dpi=80)
item_sum = df.groupby(["item"],as_index=False).agg({"sales": "sum"}).
    ↪sort_values(by="sales",ascending=False)
sns.barplot(data=item_sum,x='item',y='sales',color="green",order=item_sum.
    ↪sort_values('sales').item)
```


```
[ ]: <AxesSubplot:xlabel='item', ylabel='sales'>
```



```
[ ]: #
figure(figsize=(18, 4), dpi=80)
store_daily = df.groupby(["date","store"],as_index=False).agg({"sales":"sum"})
```


```
#
store_daily['date'] = pd.to_datetime(store_daily.date, format='%Y/%m/%d')
store_1 = store_daily[store_daily['store']==1]
ax_1 = sns.scatterplot(data=store_1,x='date',y='sales')
ax_1.set_ylabel("Sales/Store_2")
```

```
[ ]: Text(0, 0.5, 'Sales/Store_2')
```



```
[ ]: #
figure(figsize=(18, 4), dpi=80)
item_daily = df.groupby(["date", "item"], as_index=False).agg({"sales": "sum"})
#
item_daily['date'] = pd.to_datetime(item_daily.date, format='%Y/%m/%d')
item_1 = item_daily[item_daily['item']==1]
ax_2 = sns.scatterplot(data=item_1,x='date',y='sales')
ax_2.set_ylabel("Sales/Item_1")
```

```
[ ]: Text(0, 0.5, 'Sales/Item_1')
```



3

3.1 Feature Engineering

```
[ ]: def generate_timeline_features(data):
    data = data.copy()
    date_format = pd.to_datetime(data.date)
    data['year'] = date_format.dt.year
    data['month'] = date_format.dt.month
    # Q1-Q4
    data['quarter'] = date_format.dt.quarter
    #
    data['dayofweek'] = date_format.dt.dayofweek
    data['dayofyear'] = date_format.dt.dayofyear
    # 0: Winter - 1: Spring - 2: Summer - 3: Fall
    data["season"] = np.where(data.month.isin([12,1,2]), 0, 1)
    data["season"] = np.where(data.month.isin([6,7,8]), 2, data["season"])
    data["season"] = np.where(data.month.isin([9, 10, 11]), 3, data["season"])
    return data
new_df = generate_timeline_features(df)
```

```
[ ]: new_df.groupby([ "year", "month","store", "item"]).agg({"sales": ["sum",
↪ "mean", "median", "std"]}).tail(5)
```

```
[ ]:
      sales
      sum mean median std
year month store item
2018 3      10    46    0.0  NaN    NaN NaN
      47    0.0  NaN    NaN NaN
      48    0.0  NaN    NaN NaN
      49    0.0  NaN    NaN NaN
      50    0.0  NaN    NaN NaN

      #    ## Lag features
```

```
[ ]: # lags = [91, 98, 105, 112, 180, 270, 365, 546, 728]
lags = [91, 180, 365, 546]
def lag_features(df, lags):
    for lag in lags:
        value = df.groupby(["store", "item"])['sales'].transform(lambda x: x.
↪ shift(lag))
        df['sales_lag_' + str(lag)] = value
    return df
new_df= lag_features(new_df, lags)
```

4

4.1 Rolling Mean Features

```
[ ]: def roll_mean_features(df, windows):
    for window in windows:
        value = df.groupby(["store", "item"])['sales'].transform(lambda x: x.
↳ shift(1).rolling(window=window, min_periods=10, win_type="triang").mean())
        df['sales_roll_mean_' + str(window)] = value
    return df
new_df= roll_mean_features(new_df, [365, 546])
```

5

5.1 Exponentially Weighted Mean Features

```
[ ]: def ewm_features(dataframe, alphas, lags):
    for alpha in alphas:
        for lag in lags:
            dataframe['sales_ewm_alpha_' + str(alpha).replace(".", "") +
↳ "_lag_" + str(lag)] = \
                dataframe.groupby(["store", "item"])['sales'].transform(lambda
↳ x: x.shift(lag).ewm(alpha=alpha).mean())
    return dataframe
alphas = [0.9, 0.8, 0.7, 0.5]
new_df= ewm_features(new_df, alphas, lags)
new_df.tail(2)
```

```
[ ]:
      date store item sales      id year month quarter \
44998 2018-03-30    10   50   NaN 44998.0 2018     3         1
44999 2018-03-31    10   50   NaN 44999.0 2018     3         1

      dayofweek dayofyear ... sales_ewm_alpha_08_lag_365 \
44998         4         89 ...                68.550876
44999         5         90 ...                68.910175

      sales_ewm_alpha_08_lag_546 sales_ewm_alpha_07_lag_91 \
44998                82.314892                69.403475
44999                94.062978                64.221042

      sales_ewm_alpha_07_lag_180 sales_ewm_alpha_07_lag_365 \
44998                98.791375                68.602440
44999                79.337413                68.880732

      sales_ewm_alpha_07_lag_546 sales_ewm_alpha_05_lag_91 \
44998                83.485707                66.038719
44999                92.945712                64.019360
```

```

sales_ewm_alpha_05_lag_180 sales_ewm_alpha_05_lag_365 \
44998          96.603586          68.716870
44999          83.801793          68.858435

sales_ewm_alpha_05_lag_546
44998          84.936127
44999          90.968063

[2 rows x 33 columns]


```

6

6.1 Logarithmic Transformation for Sales Data

```
[ ]: new_df.hist('sales',figsize=(8,5))
```

```
[ ]: array([[<AxesSubplot:title={'center':'sales'}>]], dtype=object)
```




50

```
[ ]: #
new_df['sales_log'] = np.log1p(new_df["sales"].values)
```

```
new_df.hist('sales_log',figsize=(8,5))
```

```
[ ]: array([[<AxesSubplot:title={'center':'sales_log'}>]], dtype=object)
```



7

```
[ ]: # --One-Hot-Encoding
# store 10 50 7 4 4 6 12
# 1
# status_features = ['store', 'item', 'dayofweek', "quarter", 'month', "year", "season"]
# status_features = ['store', 'item', 'dayofweek', "season"]
#
# df_discrete = pd.get_dummies(new_df, columns=status_features)
#
# tem = df_discrete.copy()
#
# df_discrete_sorted = tem.sort_values("date").reset_index(drop = True)
# 2017
```



```

train_df = df_discrete_sorted.loc[(df_discrete_sorted["date"] < "2015-01-01"), :
    ↪]
# 2017
valid_df = df_discrete_sorted.loc[(df_discrete_sorted["date"] >= "2015-01-01")&
    ↪(df_discrete_sorted["date"] < "2015-04-01"),:]
#
cols = [col for col in train_df.columns if col not in ['date', 'id', "sales",
    ↪"year"]]
#
X_train = train_df[cols]
#
X_valid = valid_df[cols]
#
y_train = train_df['sales']
#
y_valid = valid_df['sales']

```

8 XGBoost

1. Explained Variance Score 0-1, 1 2. Mean Absolute Error(MAE) 3. R2 Score Explained Variance Score 4. Root Mean Squared Error RMSE

```

[ ]: from sklearn.metrics import explained_variance_score, mean_absolute_error,
    ↪mean_squared_error, r2_score

# SMAPE: Symmetric mean absolute percentage error (adjusted MAPE)
def smape(preds, target):
    n = len(preds)
    masked_arr = ~((preds == 0) & (target == 0))
    preds, target = preds[masked_arr], target[masked_arr]
    num = np.abs(preds-target)
    denom = np.abs(preds)+np.abs(target)
    smape_val = (200*np.sum(num/denom))/n
    return smape_val

def xgb_smape(y_pred, y_true):
    smape_val = smape(np.expm1(preds), np.expm1(y_true))
    return 'SMAPE', smape_val, False
#
xgb_model= xgb.XGBRegressor()
#
first_model= xgb_model.fit(X_train, y_train,
    eval_metric= lambda y_pred, y_true: [xgb_smape(y_pred,
    ↪y_true)])

```

```

print("VALID SMAPE:", smape(np.expm1(first_model.predict(X_valid)), np.
    ↪expm1(y_valid)))

#print("\tExplained variance:", explained_variance_score(y_valid, first_model.
    ↪predict(X_valid)))
print("\tMean absolute error (MAE):", mean_absolute_error(y_valid, first_model.
    ↪predict(X_valid)))
#print("\tRoot Mean squared error (RMSE):", np.
    ↪sqrt(mean_squared_error(y_valid, first_model.predict(X_valid))))
#print("\tR2 score:", r2_score(y_valid, first_model.predict(X_valid)))

```

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:797: UserWarning:
`eval_metric` in `fit` method is deprecated for better compatibility with
scikit-learn, use `eval_metric` in constructor or `set_params` instead.
UserWarning,

VALID SMAPE: 0.005818615765527391
Mean absolute error (MAE): 7.355786429511176e-05

/opt/conda/lib/python3.7/site-packages/ipykernel_launcher.py:22: RuntimeWarning:
overflow encountered in expm1


9 XGBoost

```

[ ]: feature_importance= pd.DataFrame(data = first_model.feature_importances_,
    index = first_model.feature_names_in_,
    columns= ['importance'])
feature_importance.sort_values('importance', ascending=False).head(25)
figure(figsize=(28, 24), dpi=80)
feature_importance.head(25).sort_values('importance').plot(kind='barh',
    ↪title='Feature Importance')
plt.legend(loc='lower right')
plt.show()

```

<Figure size 2240x1920 with 0 Axes>



10

10.1 Predication for Sales

```
[ ]: from time import time

data= df_discrete_sorted.copy()
# Sales
train = data.loc[~data.sales.isna()]
# Sales
test = data.loc[data.sales.isna()]
X_train = train[cols]
X_test = test[cols]

Y_train = train['sales']

start = time()
xgb_params= {"colsample_bytree": 0.3,
             "learning_rate": 0.1,
             "max_depth": 3,
             "n_estimators": 100,
             "verbose": 30,
             "num_boost_round": xgb_model.best_iteration}

xgbtrain_all= xgb.DMatrix(data=X_train, label=Y_train)

test_model= xgb.train(xgb_params, xgbtrain_all,
```

```

num_boost_round=xgb_model.best_iteration)

train_time = time() - start
start = time()
test_preds = test_model.predict(xgb.DMatrix(X_test))
predict_time = time()-start
test_preds

```

[15:51:32] WARNING: ../src/learner.cc:627:
Parameters: { "n_estimators", "num_boost_round", "verbose" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

```
[ ]: array([16.07758 , 28.052837, 14.076222, ..., 45.081093, 15.289815,
          41.9718  ], dtype=float32)
```

11

11.1 Submmition

Kaggle

```
[ ]: #      id sales
submission_df = test.loc[:, ['id', 'sales']]
#
submission_df['sales'] = np.expm1(test_preds)
# id
submission_df['id'] = submission_df.id.astype(int)
submission_df.to_csv('submission.csv', index=False)

```